



*Personalized Weather Forecast Video Generator using  
GenCast and AWS*

*April to September 2025*

# GenCast Forecasting Prototype

Internship Project Jacques Meyer– Cognizant UK



## Project Overview

This project explores and demonstrates the capabilities of DeepMind's GenCast weather prediction model.

→ **The goal:** to develop a prototype application that generates personalized, smartphone-format weather forecast videos based on user-selected dates and locations.

## Key Objectives

- Evaluate GenCast model variants for real-time suitability
- Explore methods to increase temporal resolution (from 12h to hourly)
- Identify best sources of real-time weather data
- Deliver a client-facing prototype showcasing GenCast's adaptability for weather-critical industries

# Deliverables



- GitHub code depo (video generation + model endpoint + UI)
- Competitor Analysis.
- Evaluation report on GenCast model performance.
- Final presentation (This slide deck)
- Project School report
- Demo video(s) for client showcase.
- Video presentation of the project and its context made by Cognizant's Marketing team

# User Case Definition Based on Model Capabilities

Based on the available variables predicted by GenCast and their respective forecasting performance, we designed a user case tailored to the strengths of the model and the needs of wind energy operations.

Selected user case :



## **Omar**

Renewable Energy Operator  
Yorkshire

### **Goals**

- Predict power output to align with commitments
- Schedule maintenance during low-wind periods

### **Pain Points**

- Forecasts are too generic for turbine-level planning
- Missed production estimates cause grid penalties

### **How GenCast Helps**

- Provides micro-local precision at high granularity
- Offers clear visuals to aid decision-making

# Explored user cases



## David

Emergency Response Coordinator  
Greater London Authority

### Goals

- Pre-position teams and supplies for severe weather
- Communicate timely warnings to the public

### Pain Points

- Uncertainty in forecasts leads to reactive decisions
- Weather data is difficult to translate into public guidance

### How GenCast Helps

- Scenario-based forecast videos support briefings
- Personalised data enables earlier planning and mobilisation



## Anita

Rail Logistics Coordinator  
UK Midlands

### Goals

- Minimise weather delays
- Optimise cargo handoffs

### Pain Points

- Corridor forecasts are not specific enough
- Updates are too slow to support re-routing

### How GenCast Helps

- Enables predictive routing across affected rail corridors
- Real-time updates support agile logistics planning



## Sarah

Port Operations Manager  
Port of Felixstowe, UK

### Goals

- Optimise container loading/unloading around weather windows
- Reduce delays from fog, wind, or storms

### Pain Points

- Forecasts are too coarse and infrequent
- Difficult to visualise data and integrate it into operations

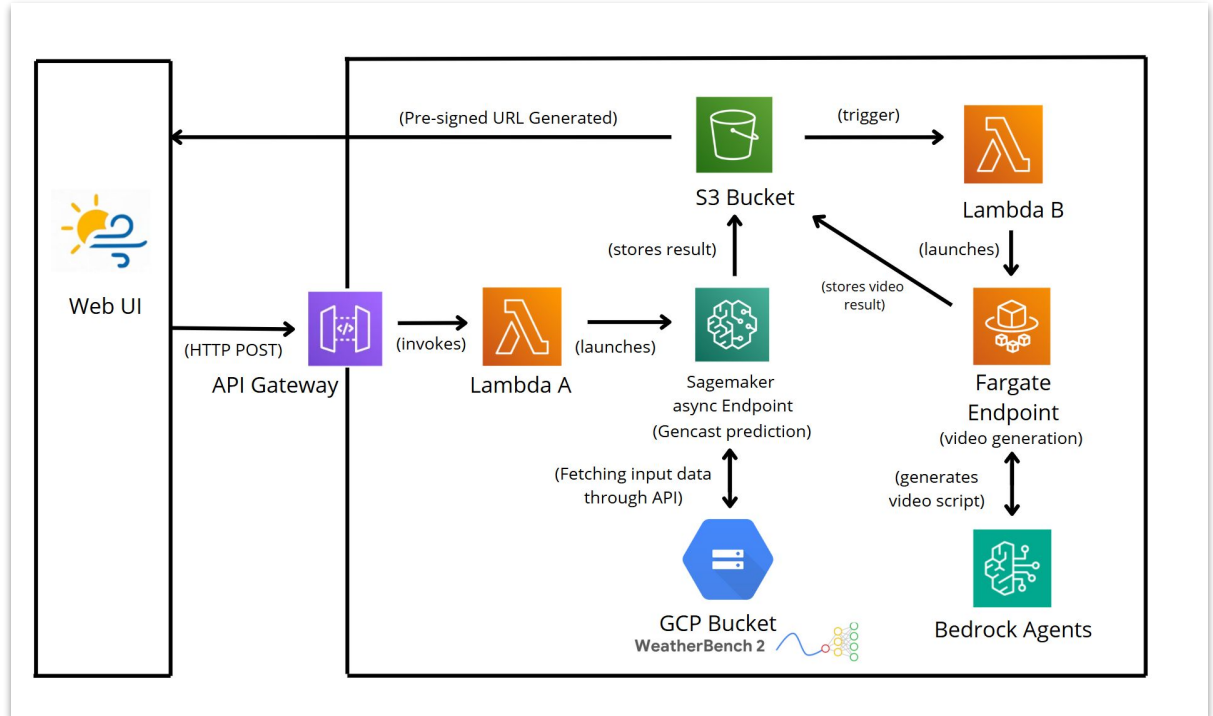
### How GenCast Helps

- Provides hourly, location-specific forecasts for port operations
- Delivers clear visuals to support scheduling and shift planning

# WeatherIQ global AWS service Architecture :

The entire project was designed  
and implemented on AWS

Global AWS service Architecture :





# User Interface

The user interface was designed using HTML and CSS, ensuring a clean and responsive layout.

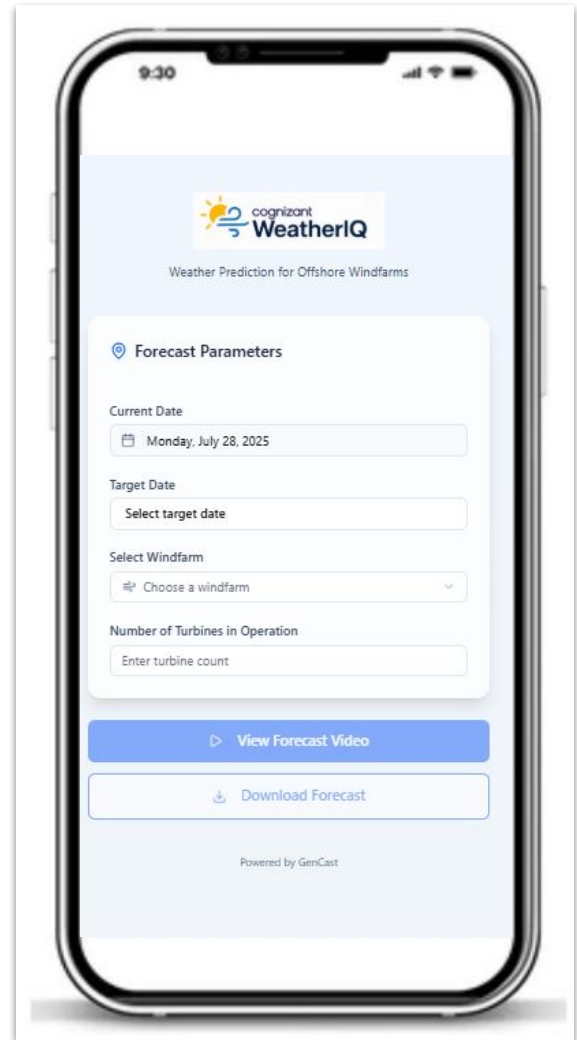
Generated using Figma's AI-powered design tool, which accelerated the development.


## User input :

- Current date (past or real time)
- Target date
- Selected windfarm
- Number of operating turbines



Weather IQ Mobile Interface.zip





# Gencast Sagemaker Endpoint

## Model Overview

17 input variables, 12 predicted variables

GenCast is a probabilistic generative AI model for weather forecasting, developed by DeepMind.

GitHub Repository <https://github.com/google-deepmind/graphcast>

- **Key Features:**

- Leverages machine learning for fast, accurate forecasts
- Outperforms the European Centre for Medium-Range Forecasts (ECMWF)'s ensemble forecast, ENS on 97.4% of 1320 targets evaluated, and better predicts extreme weather, tropical cyclones, and wind power production

- **Model Architecture**

- Autoregressive Design:  
Predicts future atmospheric states using past snapshots (e.g., from t-12h and t, forecast t+12h)
- Forecast Horizon:  
Iterative predictions up to 15 days ahead

- **Training Data**

- Dataset: ERA5 (via ECMWF & Copernicus Climate Change Service)
- Period:
  - Training: 1979–2018
  - Evaluation: 2019–present
- Resolution: 0.25° or 1° grid, 12-hour intervals

Type	Variable name	Short name	ECMWF Parameter ID	Role (accumulation period, if applicable)
Atmospheric	Geopotential	z	129	Input/Predicted
Atmospheric	Specific humidity	q	133	Input/Predicted
Atmospheric	Temperature	t	130	Input/Predicted
Atmospheric	U component of wind	u	131	Input/Predicted
Atmospheric	V component of wind	v	132	Input/Predicted
Atmospheric	Vertical velocity	w	135	Input/Predicted
Single	2 metre temperature	2t	167	Input/Predicted
Single	10 metre u wind component	10u	165	Input/Predicted
Single	10 metre v wind component	10v	166	Input/Predicted
Single	Mean sea level pressure	msl	151	Input/Predicted
Single	Sea Surface Temperature	sst	34	Input/Predicted
Single	Total precipitation	tp	228	Predicted (12h)
Static	Geopotential at surface	z	129	Input
Static	Land-sea mask	lsm	172	Input
Static	Latitude	n/a	n/a	Input
Static	Longitude	n/a	n/a	Input
Clock	Local time of day	n/a	n/a	Input
Clock	Elapsed year progress	n/a	n/a	Input

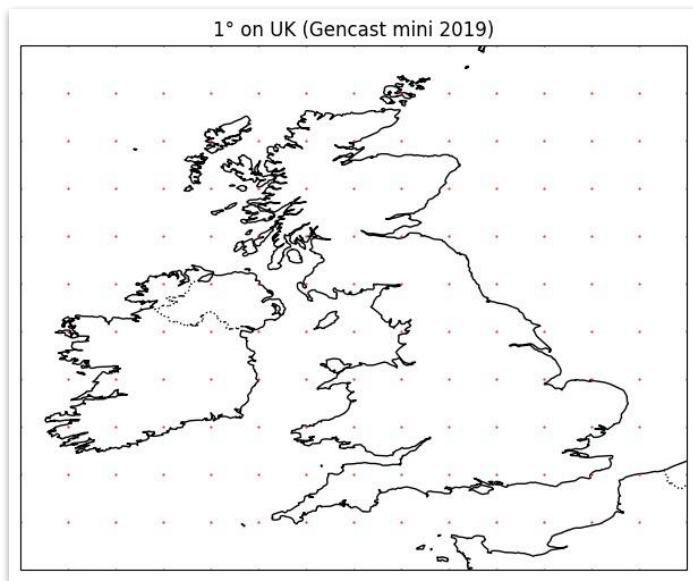




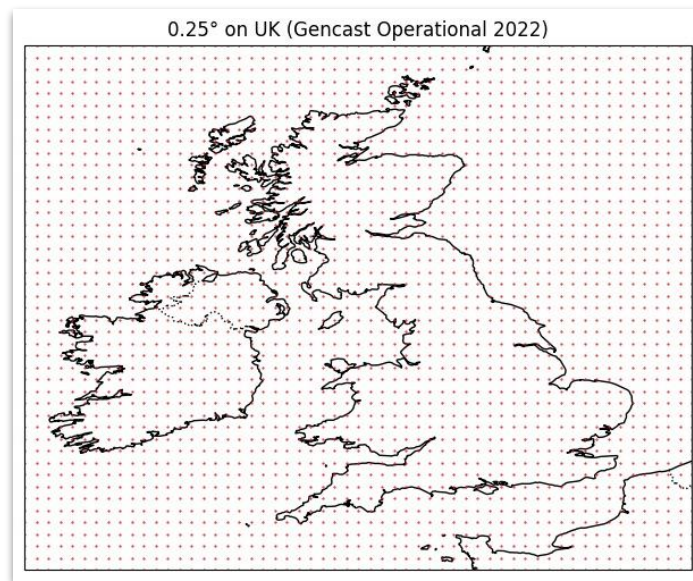
# Gencast Sagemaker Endpoint


Resolution difference between models :

[GenCast 1p0deg <2019.npz](#)



[GenCast 0p25deg <2019.npz](#)





# Gencast Sagemaker Endpoint

## Data Architecture following Weather Bench2 :

WeatherBench 2 is an open-source benchmark developed by Google Research and ECMWF to evaluate weather forecasting models—both machine learning and traditional.

It provides:

- Standardized datasets (e.g., ERA5 in cloud-optimized format)
- Evaluation tools for comparing model accuracy and realism
- Support for probabilistic and ensemble forecasts

→ It helps researchers track progress and compare models fairly across different forecasting tasks.

xarray.Dataset

~ Dimensions: (lat: 181, lon: 360, batch: 1, time: 2, level: 13)

▼ Coordinates:


lat	(lat)	float32	-90.0 -89.0 -88.0 ... 89.0 90.0	📄	☰
lon	(lon)	float32	0.0 1.0 2.0 ... 357.0 358.0 359.0	📄	☰
level	(level)	int64	50 100 150 200 ... 700 850 925 1000	📄	☰
datetime	(batch, time)	datetime64[ns]	2019-03-29 2019-03-29T12:00:00	📄	☰
time	(time)	timedelta64[ns]	-1 days +12:00:00 00:00:00	📄	☰

▼ Data variables:

land_sea_mask	(lat, lon)	float32	1.0 1.0 1.0 1.0 ... 0.0 0.0 0.0 0.0	📄	☰
geopotential_at...	(lat, lon)	float32	2.735e+04 2.735e+04 ... -0.07617	📄	☰
2m_temperature	(batch, time, lat, lon)	float32	227.9 227.9 227.9 ... 246.3 246.3	📄	☰
sea_surface_tem...	(batch, time, lat, lon)	float32	nan nan nan ... 271.5 271.5 271.5	📄	☰
mean_sea_level_...	(batch, time, lat, lon)	float32	1.003e+05 1.003e+05 ... 1.005e+05	📄	☰
10m_v_compone...	(batch, time, lat, lon)	float32	0.2381 0.2381 ... -0.4302 -0.4302	📄	☰
total_precipitatio...	(batch, time, lat, lon)	float32	0.0002059 0.0002059 ... 0.0001269	📄	☰
10m_u_compon...	(batch, time, lat, lon)	float32	0.7009 0.7009 ... -0.3383 -0.3383	📄	☰
u_component_of...	(batch, time, level, lat, lon)	float32	-0.0008879 -0.0008879 ... 0.0003166	📄	☰
specific_humidity	(batch, time, level, lat, lon)	float32	2.84e-06 2.84e-06 ... 0.0002983	📄	☰
temperature	(batch, time, level, lat, lon)	float32	215.8 215.8 215.8 ... 245.4 245.4	📄	☰
vertical_velocity	(batch, time, level, lat, lon)	float32	0.0007182 0.0007182 ... 0.02764	📄	☰
v_component_of...	(batch, time, level, lat, lon)	float32	-0.0001837 ... -0.0003536	📄	☰
geopotential	(batch, time, level, lat, lon)	float32	1.938e+05 1.938e+05 ... 382.5 382.5	📄	☰

~ Indexes: (4)

~ Attributes: (0)




# Gencast Sagemaker Endpoint

## Input Data for real time prediction :



- Source: Google Research GCP Bucket via API  
<https://github.com/google-research/arco-era5?tab=readme-ov-file#analysis-ready-data>
  - Dataset: google-research/arco-era5
  - Format: NetCDF (.nc) in WeatherBench2 structure
  - Description: Recipes for generating Analysis-Ready & Cloud-Optimized (ARCO) ERA5 datasets
- Data Type: ERA5 and preliminary ERA5T  
→ ERA5T is available with ~1 week delay (5–6 days due to ECMWF processing)



Example :

Current date	Target date	Last available data	Prediction	AWS Inference	Prediction duration
2025-08-06	2025-08-07	2025-07-31	9 days, 19 times	ml.g5.12xlarge	26.6 minutes



# Gencast Sagemaker Endpoint

## Model Inputs & Preprocessing

### Required Datasets for Inference

1. `eval_inputs`
  - NetCDF file with two input time steps (12 hours apart)
  - Contains 18 variables (12 predicted)
  - Structured by latitude, longitude, pressure levels, and time
2. `eval_targets`
  - Empty NetCDF file with same dimensions as `eval_inputs`
  - Time coordinates represent future steps to predict
  - Filled with NaN values as placeholders for model output
3. `eval_forcings`
  - Stable, non-predicted variables derived from time and location:
    - `year_progress_sin, year_progress_cos`
    - `day_progress_sin, day_progress_cos`
  - Encodes temporal context for the model

### Forcing Calculation Logic

- Compute actual datetimes by:
  - Extracting the last timestamp from `eval_inputs`
  - Adding time offsets from `eval_targets`
- Use these to calculate sinusoidal encodings for year/day progress


### Integration with GenCast Demo Code

- Uses DeepMind's `extract_inputs_targets_forcings(...)` function to:
  - Split full dataset into `eval_inputs`, `eval_targets`, and `eval_forcings`
  - Normalize time (last input time = 0)
  - Validate structure and consistency

### Adapting for Real-World Inference

- Only input data is known; targets are unknown
- Approach:
  1. Fetch ERA5 snapshots at `t-12h` and `t`
  2. Create placeholder `eval_targets` with future time steps
  3. Concatenate inputs and targets into a synthetic dataset
  4. Run `extract_inputs_targets_forcings(...)` to:
    - Extract inputs and targets
    - Auto-compute forcings
    - Ensure alignment and integrity

→ This enables robust, inference-only workflows using DeepMind's preprocessing logic.



# Gencast Sagemaker Endpoint


## AWS SageMaker instances :

Here's the comparison table for AWS SageMaker instances running the GenCast 1x0 model with 12-step prediction (6 days):

Instance Type	Cost per Hour (USD)	Inference Time (seconds)
ml.g5.8xlarge	2.448	1875.89
ml.g5.12xlarge → <b>Selected option</b>	5.672	517.64
ml.g5.24xlarge	8.144	516.96
ml.p3.2xlarge	3.060	1990.43
ml.p3.8xlarge	12.240	Too expensive
ml.p3.16xlarge	24.480	Too expensive

## Notes

- Despite all tested instances offering more than enough performance to run the GenCast 1x0 model, inference times were slower than expected based on DeepMind's published benchmarks for GPUs.
- The 0.25° model could not be launched on any available AWS SageMaker instance due to memory limitations.
- As a result, we stayed with the 1x0 model, which remains compatible for our use case.



# GenCast Sagemaker Endpoint

## Hardware Requirements & Inference Performance :

[https://github.com/google-deepmind/graphcast/blob/main/docs/cloud\\_vm\\_setup.md](https://github.com/google-deepmind/graphcast/blob/main/docs/cloud_vm_setup.md)

- **Designed for GCP**

GenCast is optimized for Google Cloud Platform (GCP) using TPUs and the JAX library  
JAX: A high-performance numerical computing library, similar to NumPy but optimized for hardware acceleration (TPUs/GPUs)

- **Attention Mechanisms**

GenCast uses two attention mechanisms:

- splash\_attention → optimized for TPUs, faster
- triblockdiag\_mha → compatible with GPUs, but slower and more memory-intensive

### Inference Time Comparison (30-step rollout, 0.25° resolution)

Model Resolution	System Memory	GPU vRAM
0.25° GenCast	~300 GB	~60 GB
1° GenCast	~24 GB	~16 GB

→ TPUs are only available on GCP, making it the preferred platform for efficient GenCast inference.

- **Memory Requirements**


→ On AWS, use a SageMaker instance that meets these specs for GPU inference.

Configuration	Inference Time	Notes
TPU + splash_attention	~8 min	Fastest setup
TPU + triblockdiag_mha	~15 min	Slower due to attention type
GPU + triblockdiag_mha (AWS)	~25 min	Slowest; only option on AWS

## Recommendation

For next Project iteration :

- Move to GCP to leverage TPUs and splash\_attention and use 0x25 model
- Code is compatible GCP, making migration straightforward



# Gencast Sagemaker Endpoint

## Sagemaker Asynchronous Endpoints :

Due to the non-negligible cost of running GenCast inferences—especially on high-performance GPU instances—we chose to host the model using SageMaker Asynchronous Endpoints.

### Why It's Cost-Efficient :

- Pay-as-you-go: Instances run only during inference, avoiding 24/7 GPU costs.
- Ideal for prototyping or low-traffic setups (e.g., one user at a time).
- Avoids the high cost of keeping powerful GPU instances constantly active.

### Downsides :

- Startup Latency:  
Each inference has a cold start delay of ~1–2 minutes, which can negatively impact user experience.
- Not Real-Time:  
Unsuitable for applications needing instant feedback.
- Extra Complexity:  
Requires managing job status, result retrieval, and S3 storage.

While asynchronous endpoints introduce a latency trade-off that affects responsiveness, they are cost-effective for prototypes. For GenCast, this setup is acceptable given the single-user context and the high cost of continuous GPU usage.



# Gencast Sagemaker Endpoint



## Reduce GenCast's 12-Hour Resolution :

### Explored Options:

1. Multiple Offset Predictions :  
Launching several predictions with time offsets and merging results.  
→ Rejected due to high resource usage and poor UX/pricing scalability.
2. Classic ML Model (ERA5-based) :  
Train a model using ERA5 data, focusing on key variables and pressure levels.  
→ Promising but requires location-specific fine-tuning. Not pursued due to time constraints.

### Selected Option:

- Simple Linear Interpolation :  
→ Easy to implement, low computational cost, and delivers good performance.








# Gencast Sagemaker Endpoint

## Output Format :

- **File Type:** NetCDF (.nc)
- **Forecast Duration :** 72 hours starting from the user-selected date (29 time)
  - Day 1: 25 interpolated time steps
  - Days 2–3: 3 time steps per day
- **Spatial Coverage :** Global ( $1^\circ$  resolution)
  - Data is not filtered to the UK or specific windfarm to maintain flexibility and reusability across different user cases.



xarray.Dataset

Dimensions: (batch: 1, time: 29, lat: 181, lon: 360, level: 13)

Coordinates:

lat	(lat)	float32	-90.0 -89.0 -88.0 ... 89.0 90.0	
lon	(lon)	float32	0.0 1.0 2.0 ... 357.0 358.0 359.0	
batch	(batch)	int64	0	
level	(level)	int64	50 100 150 200 ... 700 850 925 1000	
time	(time)	timedelta64[ns]	6 days 12:00:00 ... 9 days 12:00:00	
datetime	(time)	datetime64[ns]	...	


[29 values with dtype=datetime64[ns]]

Data variables:

10m_u_compon...	(batch, time, lat, lon)	float64	...	
10m_v_compone...	(batch, time, lat, lon)	float64	...	
2m_temperature	(batch, time, lat, lon)	float64	...	
geopotential	(batch, time, level, lat, lon)	float64	...	
mean_sea_level...	(batch, time, lat, lon)	float64	...	
sea_surface_tem...	(batch, time, lat, lon)	float64	...	
specific_humidity	(batch, time, level, lat, lon)	float64	...	
temperature	(batch, time, level, lat, lon)	float64	...	
total_precipitatio...	(batch, time, lat, lon)	float64	...	
u_component_of...	(batch, time, level, lat, lon)	float64	...	
v_component_of...	(batch, time, level, lat, lon)	float64	...	
vertical_velocity	(batch, time, level, lat, lon)	float64	...	

Indexes: (5)

Attributes: (0)



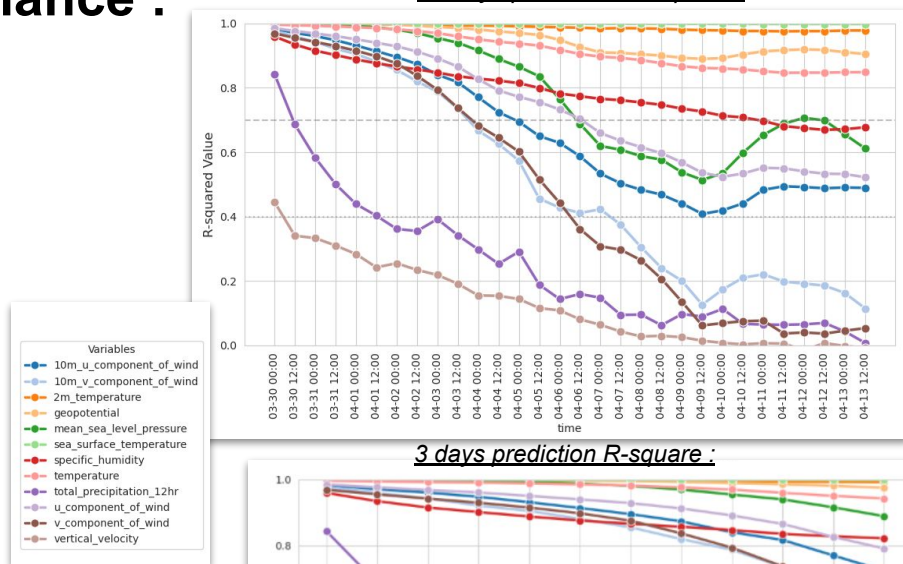
# Gencast Sagemaker Endpoint

## Variable-Specific Performance :

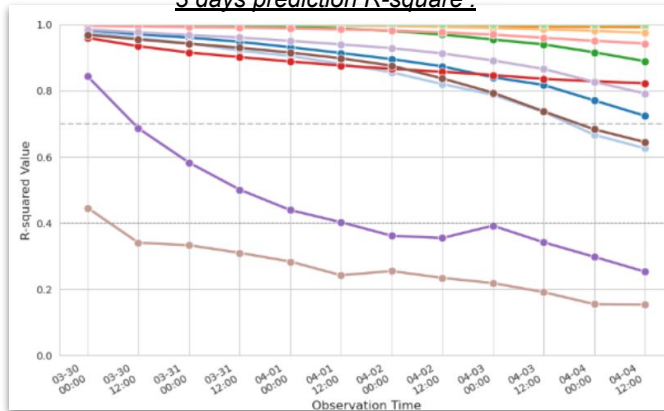
- GenCast's accuracy varies by atmospheric variable:
  - Stable variables (e.g., temperature, pressure) → easier to forecast
  - Volatile variables (e.g., wind speed, precipitation) → more challenging
- Despite these differences, GenCast consistently delivers strong results across most variables—especially when using high-quality, well-formatted input data.

→ This highlights the model's robustness and reliability for medium-range forecasting.

15 days prediction R-square :



3 days prediction R-square :





# GenCast Model: Strengths & Weaknesses

## Strengths

- **User-Friendly:** Simple to use, for technical users and good documentation.
- **Fast Inference:** Quicker predictions compared to traditional weather forecasting methods.
- **Scalable Architecture:** Can be adapted for broader use with optimization.
- **Modular Design:** Easy to integrate with other services (e.g., AWS, DeepMind APIs).
- **Single-User Optimization:** Efficient for personalized forecasts.

## Weaknesses

- **High Computational Cost:** Requires significant GPU power for high precision.
- **Limited Geographical Precision:** Accuracy drops in less-covered regions; linked to computational constraints.
- **Dependence on Real-Time Data:** Predictions degrade with outdated inputs.
- **Missing Variables:** Lacks key features like cloud coverage.
- **Latency in Real-Time Mode:** lag in input data



# Video generation Faregate Endpoint

## Forecast Video – Visual & Functional Design

Visualizations : Python + Cartopy lib for maps

### Structure & Style :


- The video presents a detailed weather forecast, inspired by traditional TV broadcasts.
- The narration walks the viewer through the day in two parts: Morning (AM) and Afternoon (PM).
- Each segment includes dynamic visualizations of:
  - Rainfall
  - Wind patterns
  - Turbine performance over time

### Temperature :

- Temperature has minimal impact on turbine performance.
- Instead, we prioritized rainfall, which supports maintenance planning.

### Performance Map :

- Shows geographical zones around the wind farm.
- Displays estimated turbine performance in each area.
- Calculated using:
  - Local wind conditions
  - The turbine's power curve
- Helps visualize how efficiently turbines would operate in different locations.





# Video generation Faregate Endpoint

## Extended Forecast Overview (3 Days)

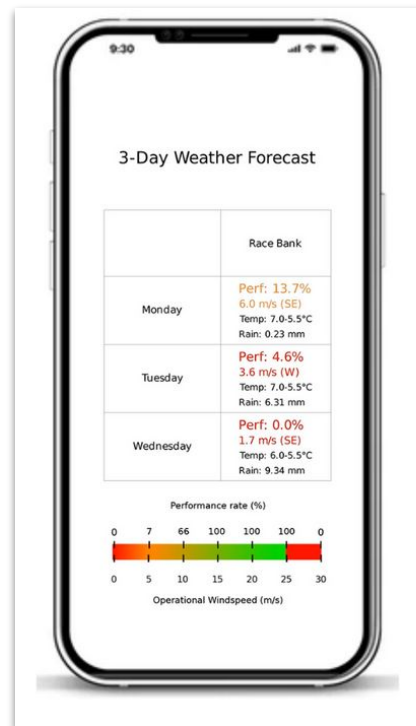
Visualizations : Python + Manim lib for animated table


### Format & Flow

Transitions to a broader overview of the next two days.

Presented in a concise tabular layout, summarizing:

- Rainfall
- Wind
- Temperature
- Turbine performance





# Video generation Faregate Endpoint

## Energy output :


Turbine documentation:

[https://www.thewindpower.net/turbine\\_en\\_20\\_siemens\\_swt-3.6-107.php](https://www.thewindpower.net/turbine_en_20_siemens_swt-3.6-107.php)

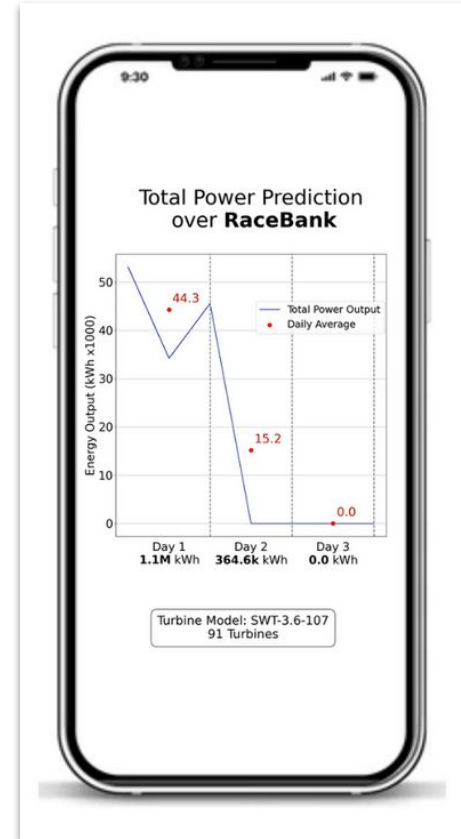
Wind speed and direction are the only variables that directly affect performance.

Incorporated the turbine's power curve to:


- Map wind speed to expected energy output for each forecasted time step
- Calculate total energy production per wind farm, based on:
  - Predicted wind conditions
  - Number of operational turbines entered by the user



*Power Curve of a SWT-3.6  
-107 Wind Turbine  
depending on the Wind  
Speed*



Visualizations :  
Python +  
matplotlib



# Video generation Faregate Endpoint



## Agentic Usage :

### Script Generation

- Bedrock agents, each responsible for a specific video section:
- All scripts were generated using Amazon's Nova Pro model, for:
  - High-quality, context-aware text generation
  - Consistent tone and structure across segments

### Audio Narration

- Used Amazon Polly with the “Amy” voice model to convert scripts into natural-sounding speech.
- Enabled seamless integration between generated content and visual storytelling



Link to the GitHub code depo :

<https://github.com/JacquesMeyerCognizant/WeatherIQ.git>





(send mail to Jacques Meyer to get access)

Link to final video demo :

[https://drive.google.com/file/d/1J8Ec\\_V2eM0IDEsaUFekqFCfu3hzz0jCk/view?usp=drive\\_link](https://drive.google.com/file/d/1J8Ec_V2eM0IDEsaUFekqFCfu3hzz0jCk/view?usp=drive_link)



# Key Objectives :

- Evaluate GenCast model variants for real-time suitability 
- Explore methods to increase temporal resolution (from 12h to hourly) 
- Real-time prediction and Identify best sources of weather data 
- Deliver a client-facing prototype showcasing GenCast's adaptability for weather-critical industries 

# Next steps and improvement :

- Move the GenCast endpoint to GCP (For pricing + 0x25 model), no adaptation to make in the code
- Finish the implementation of the global project (connecting UI + Lambdas)
- Developing new user cases

# References:

- Copernicus API :  
<https://cds.climate.copernicus.eu/datasets/reanalysis-era5-pressure-levels?tab=download>
- Deepmind forum :  
[forecasting in real time with GenCast mini demo](https://forum.deepmind.google/thread/131/forecasting-in-real-time-with-gen-cast-mini-demo) · Issue #131 · google-deepmind/graphcast
- Deepmind Era5 Bucket (recommended by Deepmind searcher) :  
[google-research/arco-era5: Recipes for reproducing Analysis-Ready & Cloud Optimized \(ARCO\) ERA5 datasets.](https://github.com/google-research/arco-era5)  
[WeatherBench 2 Data Guide — WeatherBench 2 documentation](#)
- Deepmind's public GenCast model :  
<https://github.com/google-deepmind/graphcast?tab=readme-ov-file>
- Deepmind's GenCast research paper :  
<https://www.science.org/doi/10.1126/science.adf2336>
- Wind Turbine information :  
[https://www.thewindpower.net/turbine\\_en\\_20\\_siemens\\_swt-3.6-107.php](https://www.thewindpower.net/turbine_en_20_siemens_swt-3.6-107.php)